

CSci 5105

Introduction to Distributed Systems

Naming

Today

- Naming continued
- Chapter 5 TVS, Active Names paper

Attribute-Based Naming

- Name is attribute-value pairs
- Sometimes called directory services vs. naming services

(color:blue), (size:large), ...

Example: LDAP

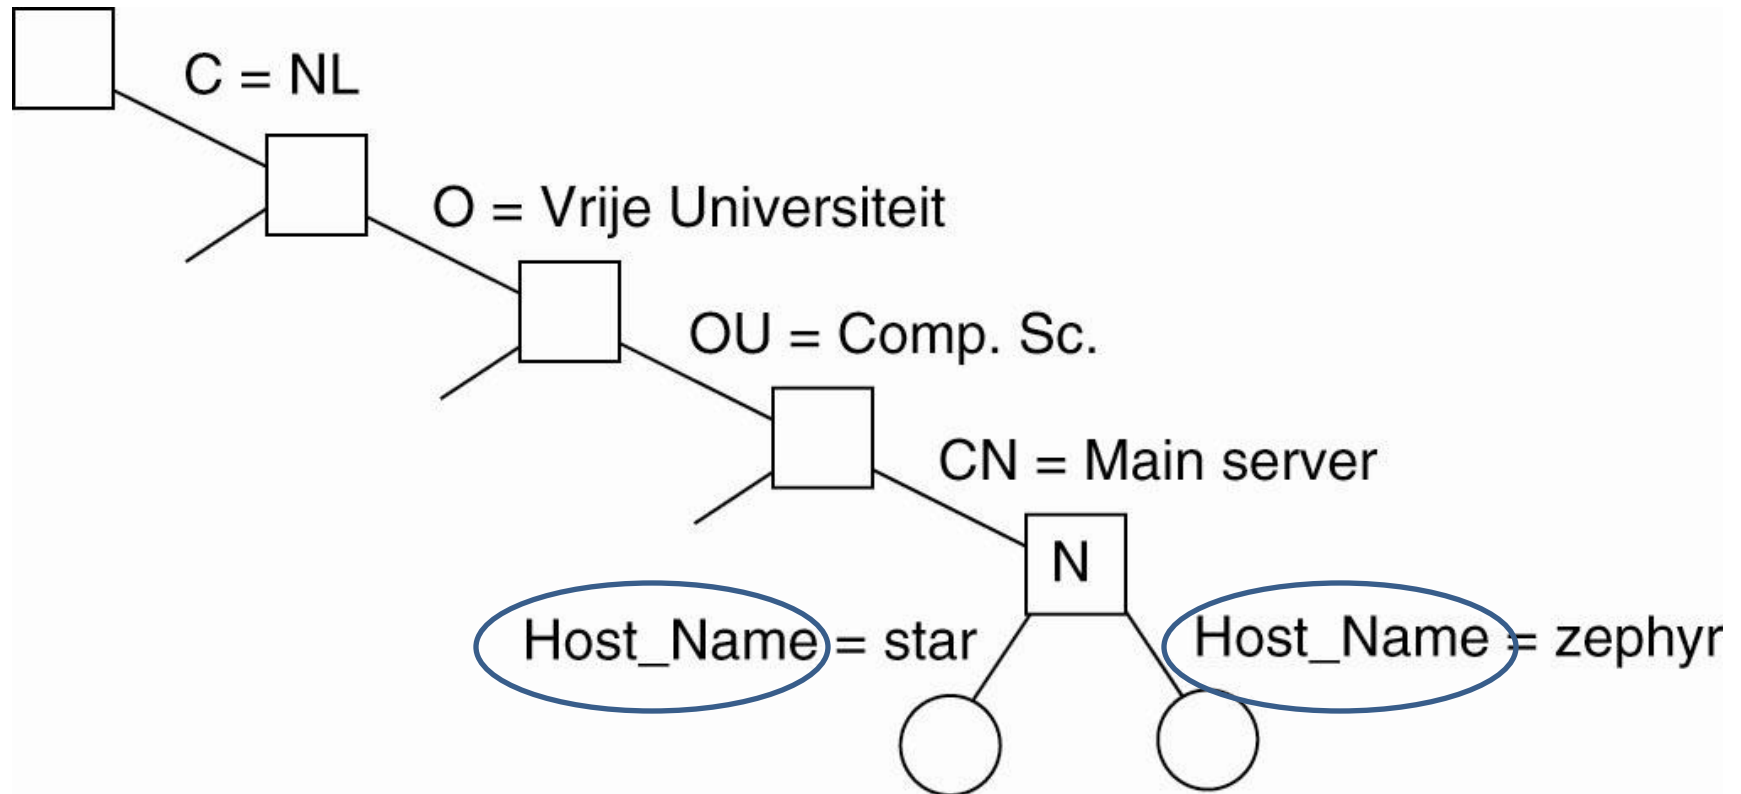
- Lightweight directory access protocol

Attribute	Abbr.	Value
Country	C	NL
Locality	L	Amsterdam
Organization	O	Vrije Universiteit
OrganizationalUnit	OU	Comp. Sc.
CommonName	CN	Main server
Mail_Servers	—	137.37.20.3, 130.37.24.6, 137.37.20.10
FTP_Server	—	130.37.20.20
WWW_Server	—	130.37.20.20

- Name is: **/C=NL/O=Vrije Universiteit/O=Comp. Sc.**


Hierarchical Implementations: LDAP

- Directory Tree



Decentralized: DHT

```
description {  
  type = book  
  description {  
    author = Tolkien  
    title = LOTR  
  }  
  genre = fantasy  
}
```



H1: Hash (type-book)

H2: Hash (type-book-author)

...

Challenge: ranges, or + and

Naming: Active Names

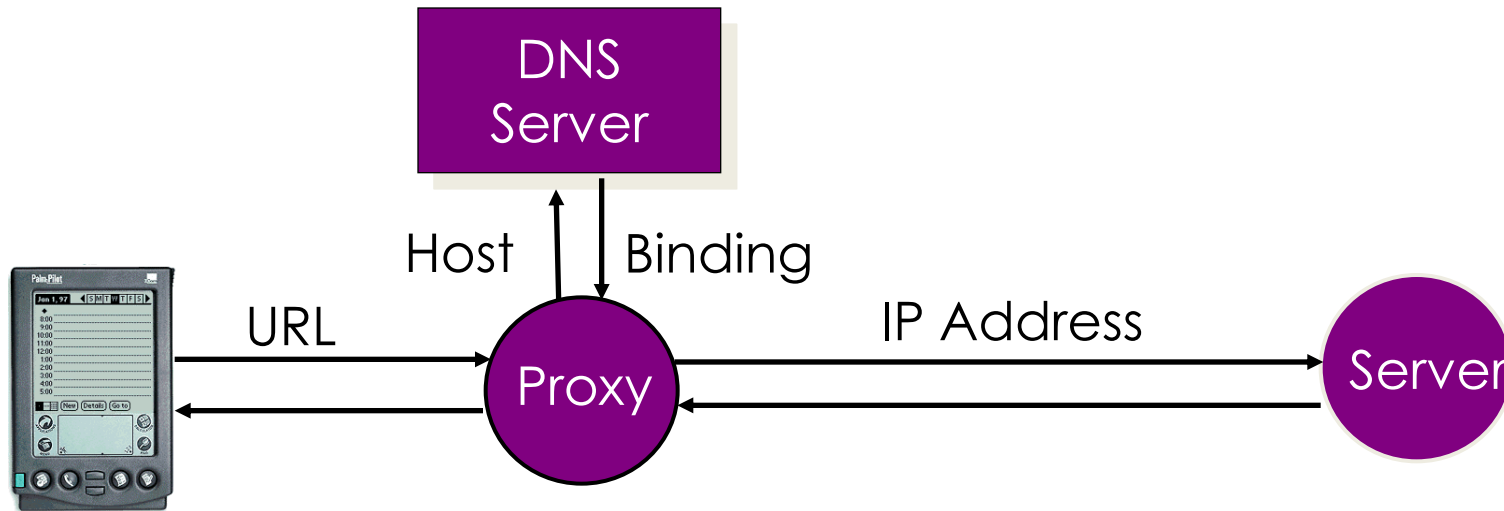
Active Names: Flexible Location and Transport of
Wide-Area Resources

Vahdat, Dahlin, Anderson, Aggarwal

The Problem

- Accessing remote resources and services is limited by rigid naming schemes
 - need a way to insert flexible (i.e. customizable) services between clients and servers
 - current solutions: done inside the network elements or deferred to application
 - either client-side or strictly server-side

Traditional Internet Naming



Traditional Model

- *Static* name -> IP address binding
- Naming and transport separate
 - http AND www.cnn.com
- Not flexible or extensible


Motivation

- Consider Scenario: context-sensitive naming
 - User types cnn.com
 - If client is behind a modem, it gets back a b/w image
 - If client is a palm pilot, it gets a distilled image
 - If the client is in Europe, it goes to the European replica
- Combine naming and transport in one framework
- Provide flexibility and extensibility in the way wide area resources are accessed

Current Attempts to Add Flexibility to Name Binding

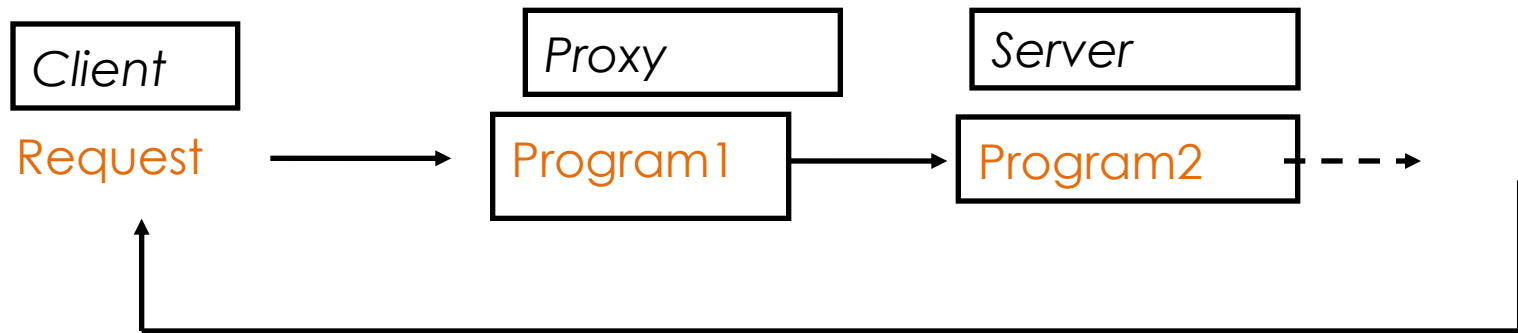
- HTTP redirect
- DNS round robin
- Cisco Local Director/Distributed Director
- Global object IDs (e.g., Globe, Legion)
- Web caches
- Mobile IP
- ... none of them are programmable

Active Names: Basic Idea



- Names resolved to mobile, secure programs
- Flexibility
- Active Names organized into hierarchical *namespaces*. A program is associated with each namespace
- Namespace programs can be changed
- Extensibility
- Active Names are connection oriented: better end-to-end semantics and performance



Programming Model



- Location independent programs
 - » Programs may run on any AN node
- Stream data model
 - » Each program operates on a data stream which is the result of the previous program
- Continuation passing style
 - » Control does not have to return to the caller program

Performance Gains


- Application customized transport protocols
- Programs are location independent. Location can be chosen to optimally utilize resources (e.g., distillation)
- Customization can be performed close to client instead of at the server (e.g., to cache dynamic content)
- 3 way RPC



Composing Services

- Delegation

- Active names organized in a hierarchy of namespaces



- Namespace programs can delegate to subordinate namespaces

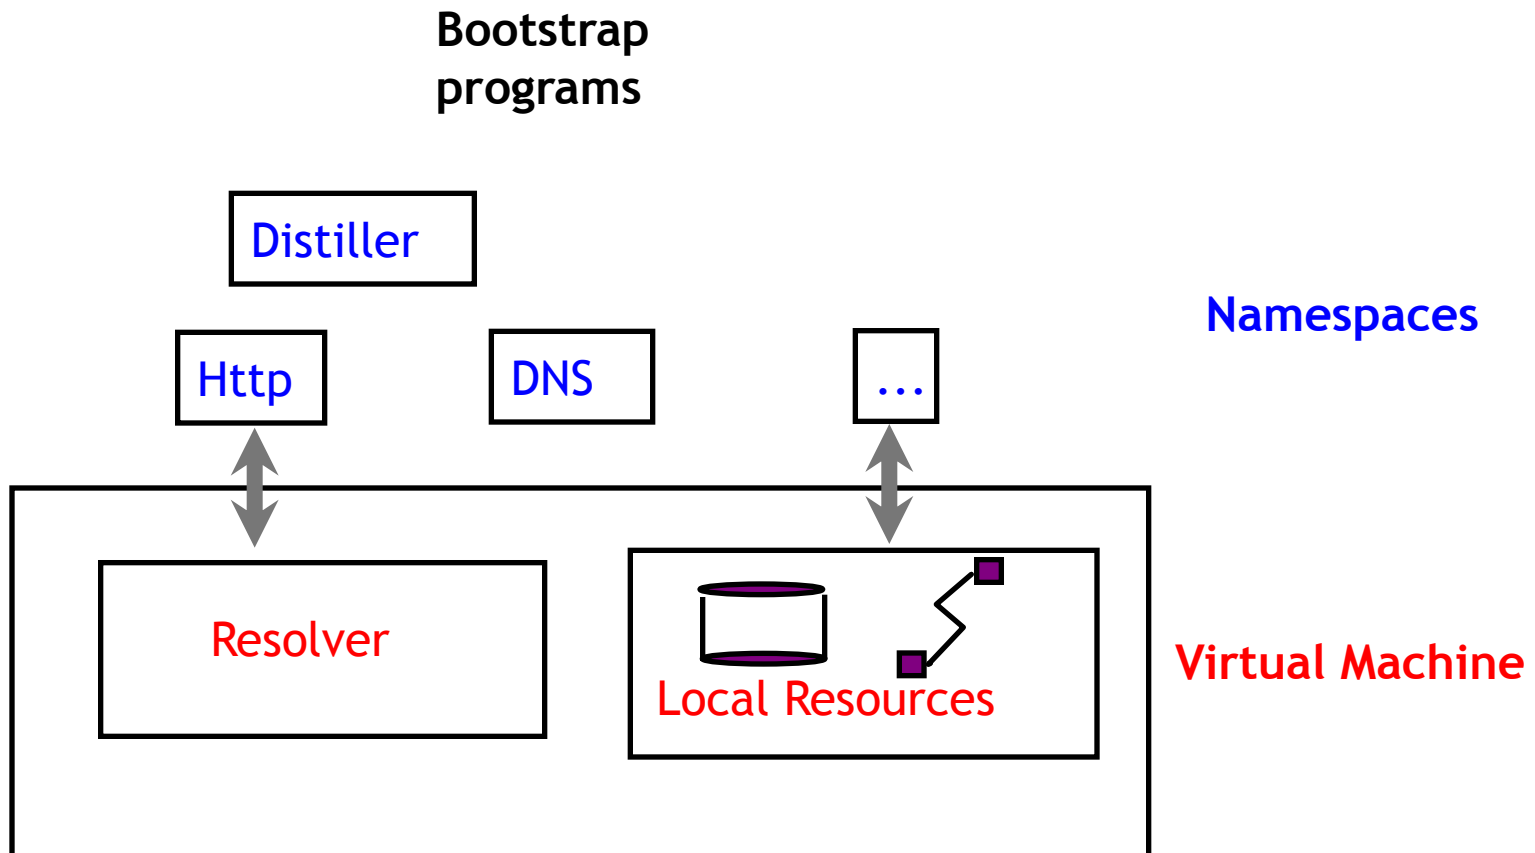
- After methods

- Continuation passing style programming
- Namespace programs bundle remaining work into “after methods” before passing control


Security

- Protection between active name programs provided by Java's type safety mechanism
- Caller passes a certificate to the callee granting it a subset of its rights
- For instance, each caller might grant its callee the right to respond to the client
- Certificates are authenticated via encryption

Active Node Architecture

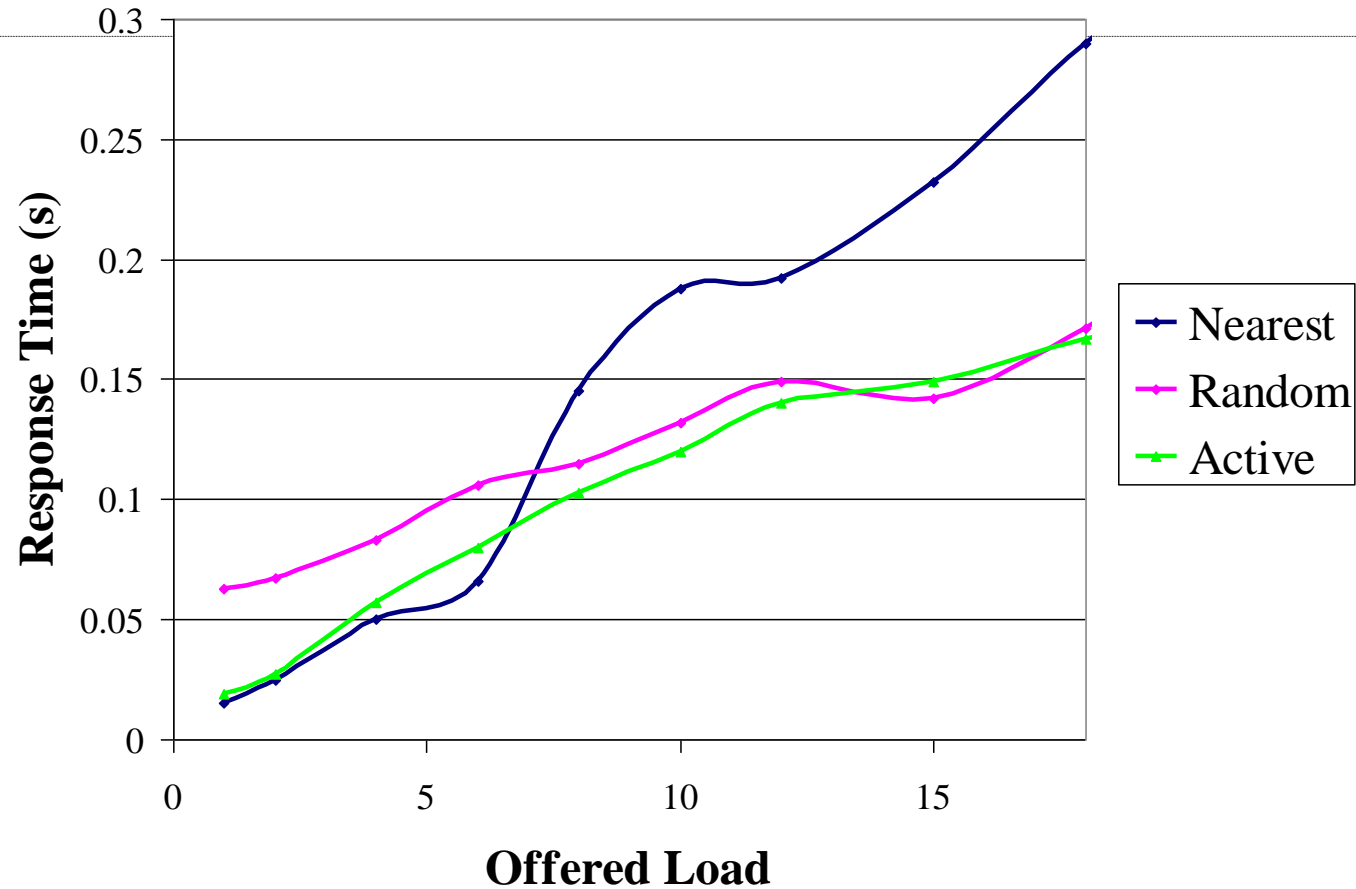


Application 1: Replica Selection



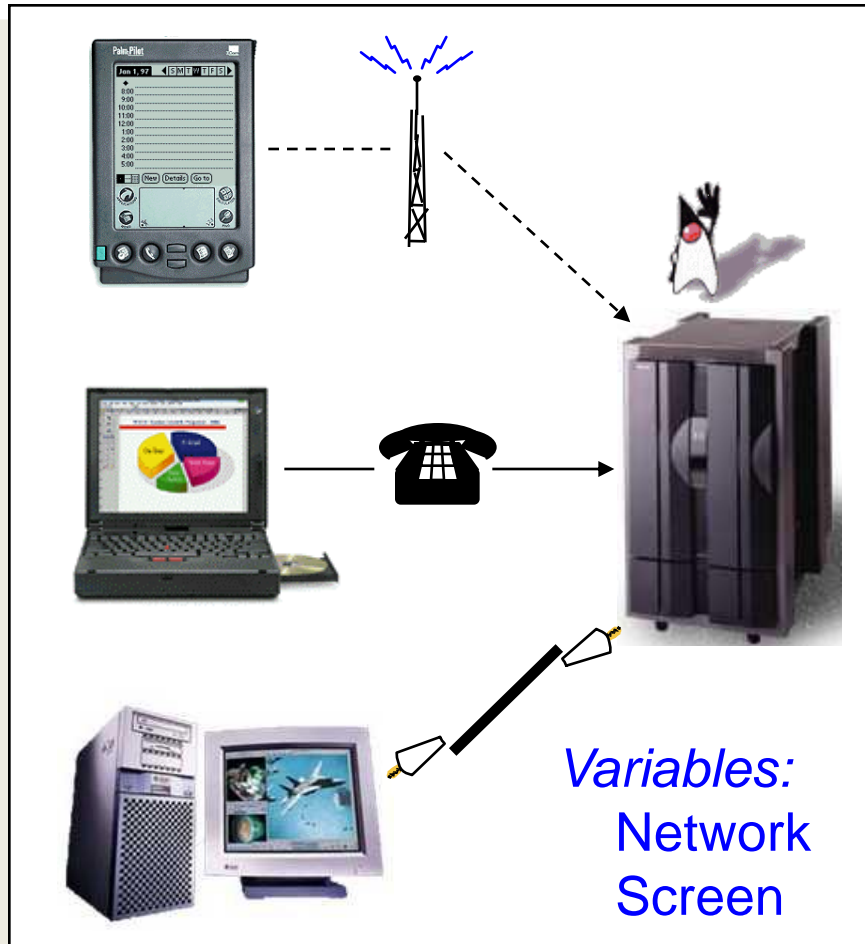
- DNS Round-Robin
 - » Randomly choose replica
 - » Avoid hot-spots
- Distributed Director
 - » Route to nearest replica
 - » Geographic locality
- Active Naming
 - » Previous performance, distance
 - » Adaptive

Replica Selection



Application 2: Mobile Distillation

Client-Specific Naming




- Clients name a single object
 - Network connection, screen
- Returned object based on client
 - Requests object, distills
- Active naming
 - Transmit name + applet
 - Flexible distillation point
 - Tradeoff computation/bandwidth
 - Support mobile clients

Application 2: Mobile Distillation

Distillation at

- Server: Saves bandwidth
- Proxy: Saves server CPU cycles
- Active: Cost estimate of both approaches



Summary

- Active name paradigm
 - Decouples name from location
 - Allows specialized processing “in the network” based on client/server conditions

Next Time

Next topic: Synchronization, Mutual Exclusion

Read Chapter 6 TVS

Have great weekend!