

CSci 5105

# Introduction to Distributed Systems

Synchronization

# Last Time


- Naming
- Fundamentals: binding, resolution, naming vs. directory service

# Time


- Why is time important in computer systems?
- Why is it hard in Distributed Systems?

# Clock Synchronization Problem

- Clocks tick at different rate, skew.
- Interrupt fires every  $K$  msec; clock updates in a register
  - may fire  $K'$



# Clock Synchronization Example



- Makefile example
- What happens?

# Solution Options

- Everyone's clock may be different
- Want notion of a single time
- Two options:


# Global time: Network Time Protocol

- Time Server B has correct time
- All other servers get in synch with it
- Server A must take what into account?


– Delay

$$\lambda = [(T_2 - T_1) + (T_4 - T_3)] / 2 \quad (\text{one-way delay})$$


$$\theta = [(T_2 - T_1) + (T_3 - T_4)] / 2 \quad (\text{offset})$$



# NTP clock filter algorithm



$$\theta = \frac{1}{2}[(T_2 - T_1) + (T_3 - T_4)]$$
$$\delta = (T_4 - T_1) - (T_3 - T_2)$$





- The most accurate offset  $\theta_0$  is measured at the lowest delay  $\delta_0$  (apex of the wedge scattergram)
- The  $\delta_0$  is estimated as the minimum of the last eight delay measurements and  $(\theta_0, \delta_0)$  becomes the peer update



# NTP in practice

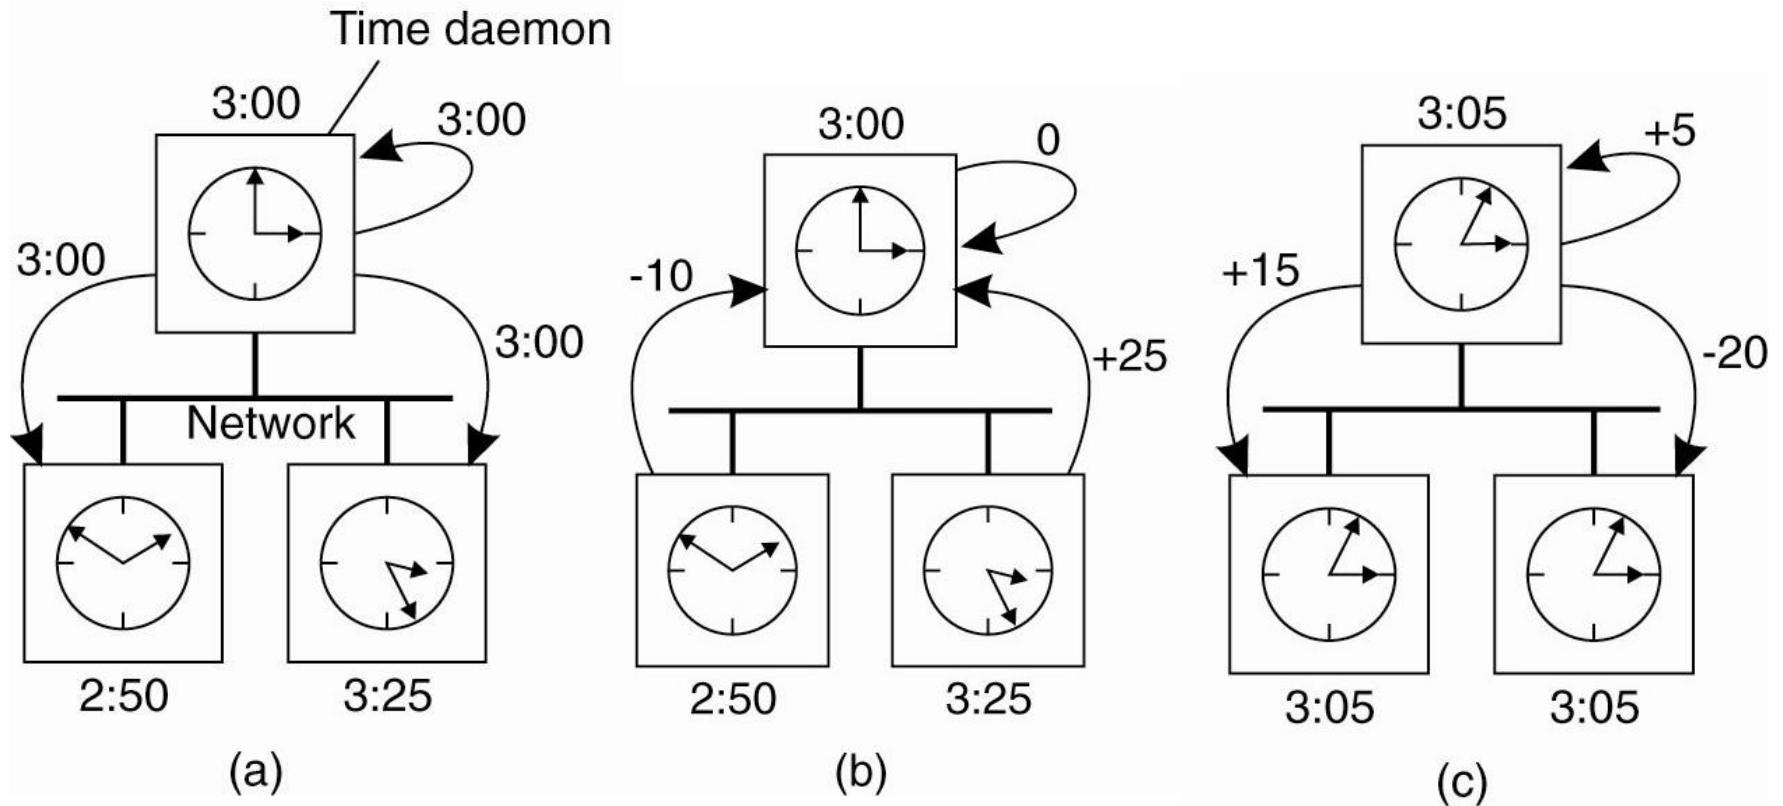
- Benefit of filtering algorithm
- Error: 724  $\mu\text{s}$  down to 192  $\mu\text{s}$



# NTP

- Challenge: if A was faster than B
  - A would have to set its clock backward
  - Set clock backward and compile .c file!
- Slowly adjust its clock backward
- 1 interrupt every 10 msec
  - + 9 msec to clock

# The Berkeley Algorithm



**No accurate global clock:** select time daemon which computes average (ignoring delay: could add back in);  
change physical clock value

# Do we need global (real) time stamps?

- Real clocks are hard to keep in synchrony
- Absolutely needed for real-time
- Maybe we can live with something weaker ...
- Makefile:
  - output.c (logical time  $L$ )
  - output.o (logical time  $L'$ )
  - $L' > L$

# Lamport's Logical Clocks

- The "happens-before" relation  $\rightarrow$  can be observed directly in two situations:
- If  $a$  and  $b$  are events in the same process, and  $a$  occurs before  $b$ , then  $a \rightarrow b$  is true.
- If  $a$  is the event of a message being sent by one process, and  $b$  is the event of the message being received by another process, then  $a \rightarrow b$


# Causal Ordering

- The ordering of events is really a partial ordering
  - Transitive  $a \rightarrow b$  and  $b \rightarrow c$  then  $a \rightarrow c$
  - Anti-symm  $a \not\rightarrow b$  and  $b \not\rightarrow a$  then  $a, b$  concurrent
  - Reflexive  $a \rightarrow a$


Total

all  $a, b$ : either  $a \rightarrow b$  or  $b \rightarrow a$

# Example



# Example



Problem?



# Quantifying Logical Clocks

- Assign clock values (#s) to events
- If  $a \rightarrow b$ 
  - $C(a) < C(b)$  for all events  $a, b$
- Note: Says nothing about order of other events


# Time Progression

Computing  $C_i$  for process  $P_i$


1. Before executing an event  $P_i$  sets  $C_i \leftarrow C_i + 1$
2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's timestamp  $ts(m)$  equal to  $C_i$
3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own local counter:

$$C_j \leftarrow \max\{C_j, ts(m)\}$$

# Lamport's Logical Clocks



(a)




(b)

# Partial vs. Total Order

- Basic lamport clocks give a partial order

- Many events happen “concurrently”
- $C(a) < C(b)$  does not imply  $a \rightarrow b$  !




- But sometimes a total order is more convenient
  - e.g., commit operations to a database
  - deposit, withdrawl
- Tie-breakers: concatenate unique PID ( $p1 > p2 > \dots$ )
  - E.g.  $C(a, P1) < C(b, P2)$  does imply  $a \rightarrow b$

# Lamport Clocks

- Cannot solve total order
  - For any two events,  $a \rightarrow b$  or  $b \rightarrow a$
- BUT can guarantee something weaker
  - Everyone acts on messages in the same order

# Using Lamport Clocks : Totally Ordered Multicasting

- Problem?




# Can solve ...

- Multicast acknowledgement
- Order events in the queue
- Execute event at top of queue in time order
  - Only if received acknowledgement to the event

# Vector Clocks

Want  $C(b) < C(a) \Rightarrow a \leftarrow b$  (and vice-versa)





# Vector Clocks


Vector clocks are constructed by letting each process  $P_i$  maintain a vector  $VC_i$  with the following two properties:

1.  $VC_i[i]$  is the number of events that have occurred so far at  $P_i$ . In other words,  $VC_i[i]$  is the local logical clock at process  $P_i$
2. If  $VC_i[j] = k$  then  $P_i$  knows that  $k$  events have occurred at  $P_j$ . It is thus  $P_i$ 's ~ knowledge of the local time at  $P_j$

# Multicast Using Vector Clocks

1. Before executing an event  $P_i$  executes  
$$VC_i[i] \leftarrow VC_i[i] + 1$$
2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's (vector) timestamp  $ts(m)$  equal to  $VC_i$
3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own vector by setting  
$$VC_j[i] \leftarrow \max\{VC_j[i], ts(m)[i]\} \text{ for each } i$$
4. If message is out-of-order, queue  $m$  !  
there exists a  $k, k \neq i$ , for which  $VC_j[k] \neq m[k]$

# Enforcement



# Next Time

Next topic: Mutual Exclusion

Read Chapter 6 TVS