

CSci 5105

Introduction to Distributed Systems

Consistency

Consistency


- **Case 1:** Replication + Updates
- Data are replicated for availability
- Data are replicated for performance
 - Scaling in numbers
 - Scaling in geographical area
- Performance/Availability gain is not free

Consistency

- Replication is not the full story
- **Case 2:** Single data accessed by many concurrent processes/users
- Two types of consistency
 - data centric, client centric

Data-centric Consistency Models


- Data is replicated at K servers



Continuous Consistency


- *Conit*: consistency unit
- value: numerical deviation
- staleness: order deviation

Replica A



Vector clock A = (15, 5)
Order deviation = 3
Numerical deviation = (1, 5)


Replica B




Vector clock B = (0, 11)
Order deviation = 2
Numerical deviation = (3, 6)

Granularity of Conit

- Policy: deviation at most one update
- Tradeoffs?



Small: (-) more overhead, (+) more accurate
Large: (-) false sharing, (+) less overhead




Qualitative Consistency

- Continuous consistency is quantitative
- Set a bound on order or numerical deviation
 - Nice! can be measured, but how to set it?
- Now, let's look at qualitative consistency based on ordering

Data-centric Consistency Models

- Data is replicated at K servers
 - Each process assumed to access one of the data-store replicas



Sequential Consistency

- Concurrent access to replicated data
 - Each process sees their operations in the sequential order
 - All processes see (via read) same global interleaving of writes

Sequential Consistency Examples

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)b	R(x)a

(a)

P3 and P4 see W(x)b then W(x)a

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)a	R(x)b

(b)

Violated - why?

Very expensive

- Why?

Sequential Consistency

- There are many valid sequences and many invalid ones

Process P1	Process P2	Process P3
$x \leftarrow 1;$ $\text{print}(y, z);$	$y \leftarrow 1;$ $\text{print}(x, z);$	$z \leftarrow 1;$ $\text{print}(x, y);$

Sequential Consistency

```
x ← 1;  
print(y, z);  
y ← 1;  
print(x, z);  
z ← 1;  
print(x, y);
```

```
x ← 1;  
y ← 1;  
print(x, z);  
print(y, z);  
z ← 1;  
print(x, y);
```

```
y ← 1;  
z ← 1;  
print(x, y);  
print(x, z);  
x ← 1;  
print(y, z);
```

```
y ← 1;  
x ← 1;  
z ← 1;  
print(x, z);  
print(y, z);  
print(x, y);
```

All valid as they represent legal interleavings
provided all processes see the same interleaving!

E.g. assignment must precede the `print` for each P_i

Stronger?

- Global/Total order
 - Time-stamp every operation
 - There can be only 1 order!

Strongest

- Strict
 - Every action occurs in the order prescribed by a real global clock
 - read sees most recent write based on actual clock

Serializability

- Stronger than sequentially consistent
- From database transactions
 - not just between reads and writes
- Execution appears to be some serial uninterrupted order of *each* P_i
- Example: $P_1P_2P_3$ or $P_2P_3P_1$

Weaker ~ Causal Consistency

- Causal ~ happened before; earlier write may have influenced later write

P1:	W(x)a		W(x)c	
P2:	R(x)a	W(x)b		
P3:	R(x)a		R(x)c	R(x)b
P4:	R(x)a		R(x)b	R(x)c

- Writes that are potentially causally related must be seen by all processes in the same order
- Concurrent writes may be seen in a different order on different machines

Causal Consistency

P1:	W(x)a		W(x)c	
P2:		R(x)a	W(x)b	
P3:		R(x)a		R(x)c
P4:		R(x)a		R(x)b

- Causally consistent
 - $W(x)a \Rightarrow W(x)c$ (everyone must see a then c)
 - $W(x)b, W(x)c$ concurrent
- It is also sequentially consistent?
 - No, all writes not seen in same order

Causal Consistency

P1: W(x)a

P2: R(x)a W(x)b

P3: R(x)b R(x)a

P4: R(x)a R(x)b

(a)

NOT OK

P1: W(x)a

P2: W(x)b


P3: R(x)b R(x)a

P4: R(x)a R(x)b

(b)

Why OK?

One More: Entry consistency



Enter a CS, data “owner” must transfer most up-to-date copy

Client-Centric

- Data-centric consistency based on ordering of events at the data
- Client-centric models are based on what a client (or a process) sees
- No simultaneous updates; single client, but client may move to another replica
- Notation: WS is the set of write operations seen at a replica

Eventual Consistency


- Context: data replicas
- An update eventually propagates to all replicas; DNS binding update
- When is this ok from client viewpoint?
 - clients access the same replica
- Not ok?

Monotonic Reads

- If a process reads the value of a data item x any successive read operation on x by that process will always return that same value or a more recent value **AND**
- If a process reads ... at **any** replica; that replica must see every prior write that was read by the process

Monotonic Reads

think of X as a
bulletin board




yes

(a)

see at post at L1
also see it later
at L2

Time along x axis




no

(b)

Li are replicas; single client or “process”

Monotonic Writes


- A write operation by a process on a data item x is completed before any successive write operation on x by the same process



(a)

yes

post at L1
post reply at L2




(b)

no

reply must follow
original post


Read Your Writes

- The effect of a write operation by a process on data item x will always be seen by a successive read operation on x by the same process
 - write completes before read



yes


post at L1 seen
later at L2



no

Writes Follow Reads


- A write operation by a process on a data item x following a previous read operation on x by the same process is guaranteed to take place on the same or a more recent value of x that was read
 - Read must complete before the write



(a)

yes

read at L1 seen
later when posting
at L2; follow up



(b)

no

Implementation

- How can we implement them?
 - Blocking
 - Vector time stamps

Next Time

Next topic: Replication

Read Chapter 7 TVS