

CSci 5105

Introduction to Distributed Systems

Replication

Today

- Replication
- Implementation of consistency protocols
- Chapter 7 TVS

Replication \Leftrightarrow Consistency

- Data are replicated for availability
- Data are replicated for performance
 - Scaling in numbers: throughput
 - Scaling in geographical area: latency

When to replicate?


- Proactive
 - a-priori
- Reactive
 - as needed
- Tradeoffs?

Replica-Server Placement

- How many server replicas?
- Where should they be placed in the network?
 - Not often studied
- Node (dot) is an interested data client
 - Each cell is handled by a replica
 - Place replica in K most dense cells; K is given

Too small => too many replicas

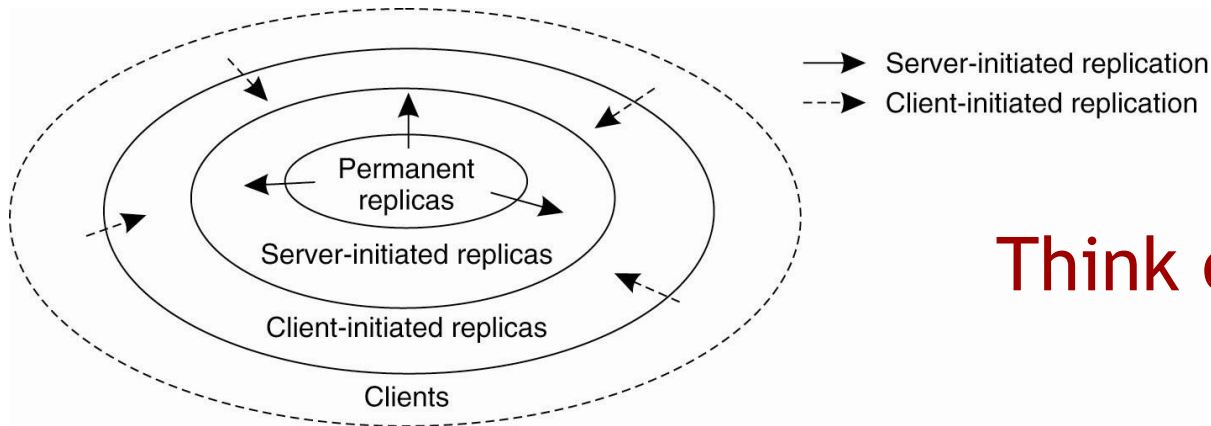
Too large => too few replicas



If same size cell requirement

Content Replication and Placement

- Types of content placement



Think of web content

- Permanent
 - distributed data-store (across servers~local domain)
 - geo-distributed mirrors
- Client-initiated: caching

Server-Initiated Replicas

- Simple scheme (how many + ~ placement)
 - Server counts client access to file F
 - Assume clients are sent to nearest server P which in turn routes to server S holding F
 - $\#Req(F, S) > rep\text{-}threshold$, replicate F
 - $\#Req(F, S) < del\text{-}threshold$, delete F (unless last one)
 - $rep\text{-}threshold$ is usually $> del\text{-}threshold$
 - If $> \frac{1}{2}$ of requests arrive to server Q from a particular server P , migrate F from Q to P
 - Why?

Keeping replicas in synch:
update propagation

How to propagate data updates between replicas?

Consistency tells us what to propagate, but
different options for *how*

1. Propagate only a notification of an update
2. Transfer data from one copy to another
3. Propagate the update *operation* to other
copies: **Active replication**

Who initiates update propagation?

- Replicas include client caches
- Push: server-based
- Pull: client-based
- Tradeoffs

Issue	Push-based	Pull-based
State at server	List of client replicas and caches	None
Messages sent	Update (and possibly fetch update later)	Poll and update
Response time at client	Immediate (or fetch-update time)	Fetch-update time

if update
notification

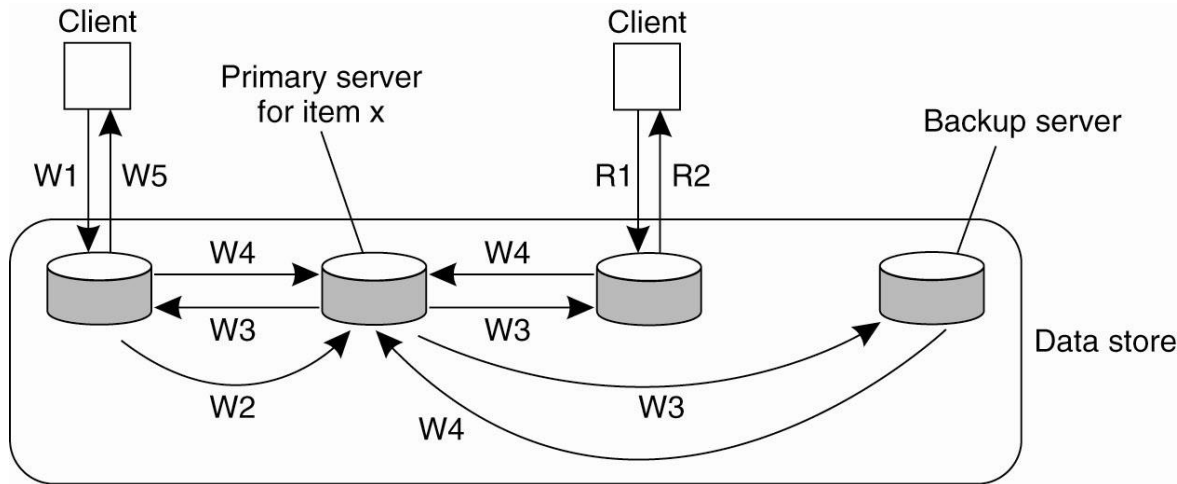
Leases

- Timed control: optimizations
- Promise server will push updates for a specified time
 - `hockey.com` **only wants** `nbcsports.com` updates until the hockey gold medal game finishes
 - after this client must poll
 - allows switching between push and pull
- Age-based lease
 - data is guaranteed to be valid for time K
 - after K , can poll

Implementation

- How are consistency models actually implemented?
- Protocol and architecture

Remote-Write Protocols: Primary-Backup



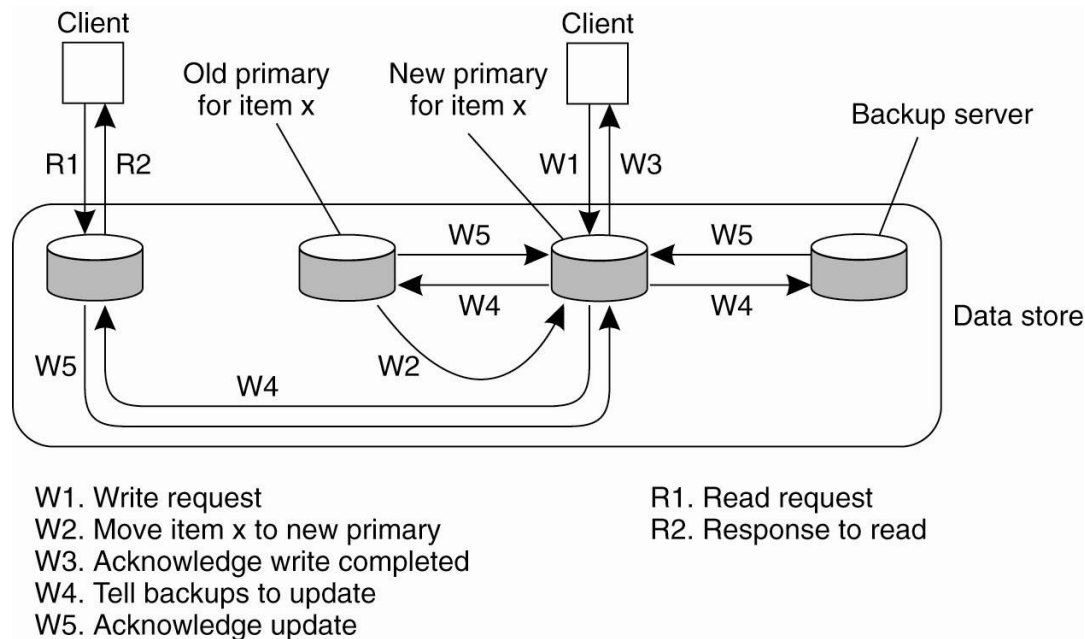
W1. Write request
W2. Forward request to primary
W3. Tell backups to update
W4. Acknowledge update
W5. Acknowledge write completed

R1. Read request
R2. Response to read

- Good for sequential consistency
- Weakness?

Local-Write Protocols

- Primary migrates to current writer
- Advantages? Disadvantages?




Replicated-Write Protocols

- These schemes write to a **single** primary
 - Global primary
 - Changing primary
 - Problems?
- Voting scheme
 - More general
 - Writes to any replica

Quorum-Based Protocols

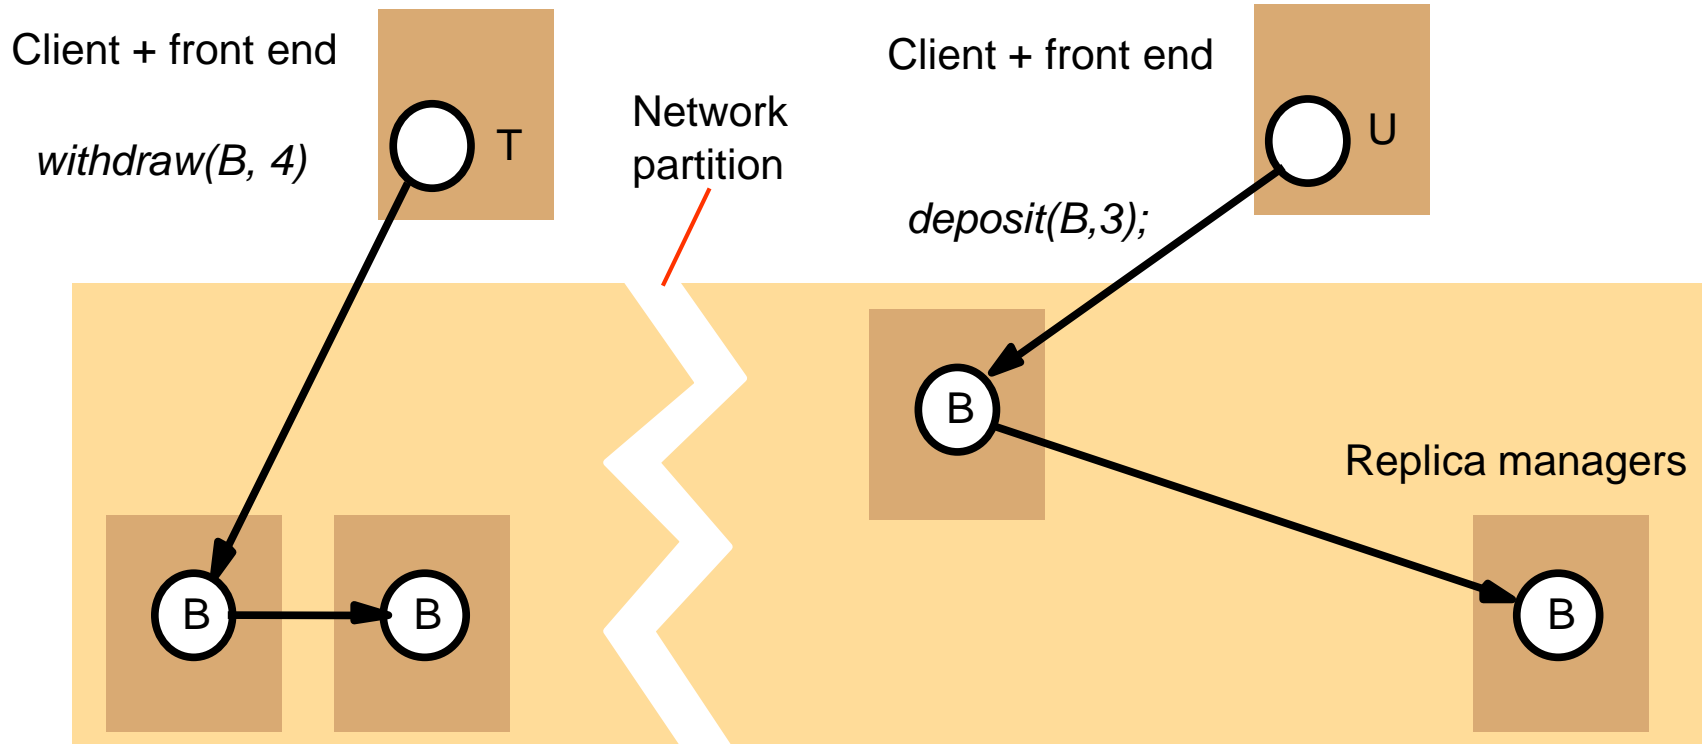
- Voting: N servers
 1. $N_r + N_w > N$ (read-write conflicts)
 2. $N_w > N/2$ (write-write conflicts)



Client-centric consistency

- Many client replicas
 - assume one server or servers are consistent
 - cache coherence ~ DFS
- Read-only caches: updates only at server
 - server-directed vs. client-directed coherence
 - server invalidates vs. client polling
- Read/write cache
 - write-through: write locally and to server
 - write-back: write locally, delay to server
 - same coherence choices

Network Partitions



- Uh Oh: how can we achieve any kind of consistency?

Next Time

Next topic: Fault Tolerance

Read Chapter 8 TVS

Have a great weekend!