

CSci 5105

# Introduction to Distributed Systems

Fault Tolerance

# Last Time

- Replication and Consistency

# Today

- Fault tolerance
- Chapter 8 TVS

# Fault Tolerance Basics

- **Availability**
  - short time horizon
  - e.g down 1 msec every hour => 99.9999 avail
- **Reliability**
  - over longer time horizon
  - e.g. but not that reliable, no job can run > 1 hr
- **Safety**: temporary failure # catastrophe
- **Maintainability**: ease of repair

Brewer Avail

# More Definition

- Fail: cannot meet promises
- Error: system state may  $\Rightarrow$  failure
- Fault: cause of an error
- Tolerate faults  $\Rightarrow$  operate correctly
- Fault types
  - Transient, intermittent, permanent

# Failure Models

Type of failure	Description
Crash failure	A server halts, but is working correctly until it halts
Omission failure <i>Receive omission</i> <i>Send omission</i>	A server fails to respond to incoming requests A server fails to receive incoming messages A server fails to send messages
Timing failure	A server's response lies outside the specified time interval
Response failure <i>Value failure</i> <i>State transition failure</i>	A server's response is incorrect The value of the response is wrong The server deviates from the correct flow of control
Arbitrary failure	A server may produce arbitrary responses at arbitrary times

byzantine

# Failure Types

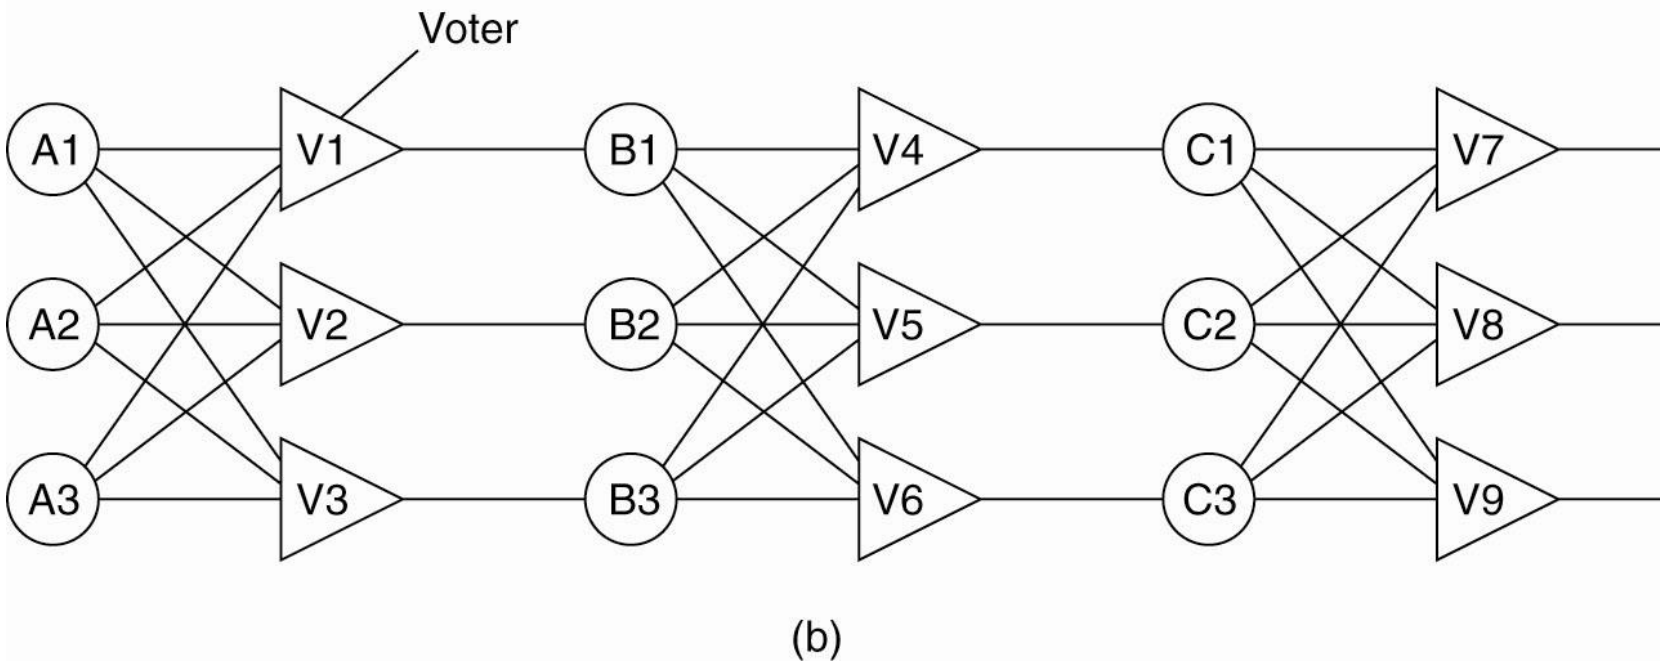

- fail-stop ~ crash failure
  - failed process stops producing output; easily detected as failed without ambiguity
  - machine on my local network
- fail-silent
  - failure not so obvious: really slow or failed?
  - remote communicating process
- fail-safe
  - arbitrary failures that are recognized as such



# RPC Failures

1. The client is unable to locate the server
    - raise exception
  2. The req. message from the client to the server is lost
  3. The server crashes after receiving a request
  4. The reply message to the client is lost
- 2-4 Detect via time-out; take action (retransmit or not)
- The client crashes after sending a request
- orphan - problem?


# Failure Masking by Redundancy



Classic TMR: throwing hardware at the problem  
Assumptions?

# Process Failures

- Process replication or groups
- Need to have group consensus
- Group can change: group management becomes key



- Compare?

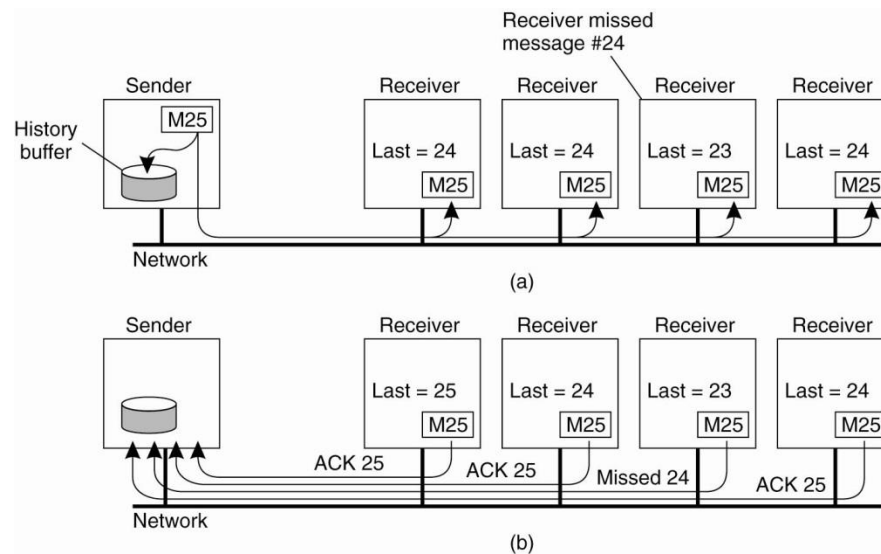
~ primary backup

# Failure Masking + Replication

- General groups
  - K fault tolerant (K failaures)
    - fail-stop/fail-silent =>
    - byzantine failures =>

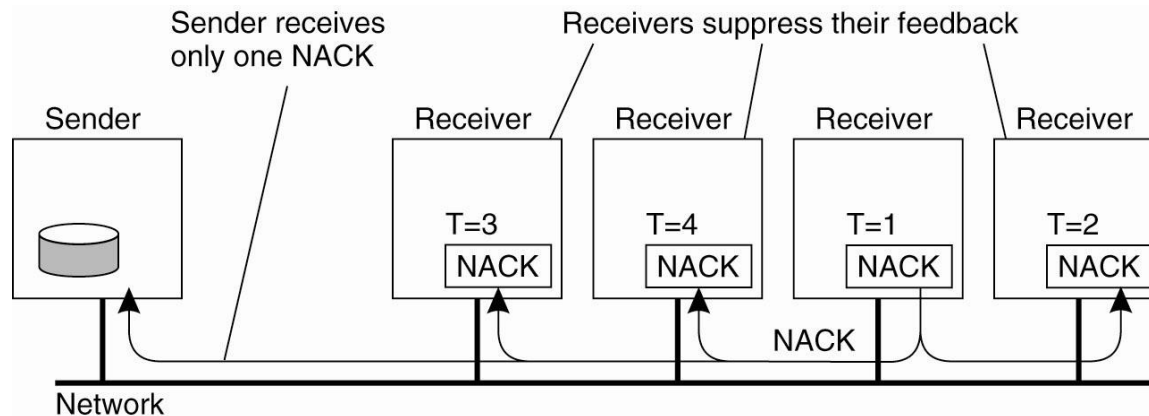
# Agreement in Faulty Systems

- Examples
  - voting, leader election, multicast
- Reliable multicast
  - group is fixed
  - failure reported via feedback



# Feedback Control

- Missing a message can unicast or multicast
- K missing: K unicasts or multicasts
- Latter: nice optimization
  - delay a little before requesting retransmission
  - another node may do it
  - So maybe 1 retransmitted multicast will suffice



# Atomic Multicast


- Reliable multicast and ordering
- Everyone sees same message order or none
- Eg. Consistency => DB updates
- Problem: group members come and go
- Agree who is in the group
  - View synchronous

# Virtual Synchrony

- Group view
  - When message  $M$  is sent; everyone agrees who is in the group
  - If group state changes during  $M$ 
    - $M$  delivered to all before group change or to none
- This is known as virtual synchrony



# Virtual Synchrony



# Multicast Message Ordering

- Unordered multicasts
- FIFO-ordered multicasts
  - Easy: issue message in sequence order
- Causally-ordered multicasts
  - Harder: need vector time-stamps
- Totally-ordered multicasts
  - Need a global sequencer
  - Each multicast message is given a global #: 1, 2, 3, ...

# Message Ordering

## Process P1

sends m1  
sends m2

## Process P2

receives m1  
receives m2

## Process P3

receives m2  
receives m1

## Process P1

sends m1  
sends m2

## Process P2

receives m1  
receives m3  
receives m2  
receives m4

## Process P3

receives m3  
receives m1  
receives m2  
receives m4

## Process P4

sends m3  
sends m4


- What ordering do these satisfy?

# Two-Phase Commit (2PC)

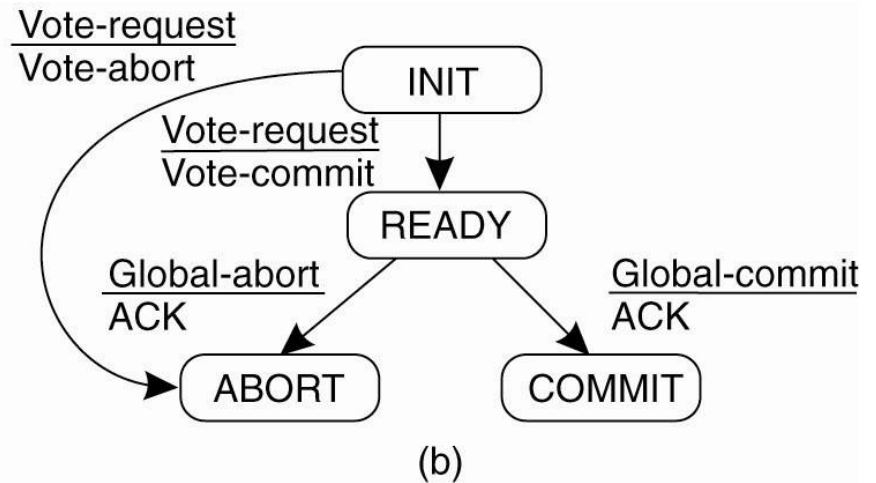
- Send message and have everyone either act on message or not
- Typical action: commit a transaction
- Multi-step
  - Vote-request
  - Vote-commit or vote-abort
  - Global-commit or global-abort
- Impressions?

# Two-Phase Commit (2PC)

Coordinator



participant



- Distributed commit - all or none

# What about failure?

- Coordinator failure
- Node P in READY state and times out
- Asks node Q

State of Q	Action by P
COMMIT	Make transition to COMMIT
ABORT	Make transition to ABORT
INIT	Make transition to ABORT
READY	Contact another participant

# 2PC Failure/Recovery

- Nodes fail and may recover
- Use logging

## Actions by coordinator:

```
write START_2PC to local log;
multicast VOTE_REQUEST to all participants;
while not all votes have been collected {
    wait for any incoming vote;
    if timeout {
        ...      write GLOBAL_ABORT to local log;
                  multicast GLOBAL_ABORT to all participants;
                  exit;
    }
    record vote;
}
```

# 2PC Failure/Recovery (cont'd)

...

```
if all participants sent VOTE_COMMIT and coordinator votes COMMIT {  
    write GLOBAL_COMMIT to local log;  
    multicast GLOBAL_COMMIT to all participants;  
} else {  
    write GLOBAL_ABORT to local log;  
    multicast GLOBAL_ABORT to all participants;  
}
```



# 2PC: Participant recovery

## actions by participant:

```
write INIT to local log;
wait for VOTE_REQUEST from coordinator;
if timeout {
    write VOTE_ABORT to local log;
    exit;
}
if participant votes COMMIT {
    write VOTE_COMMIT to local log;
    send VOTE_COMMIT to coordinator;
    wait for DECISION from coordinator;
    if timeout {
        multicast DECISION_REQUEST to other participants;
        wait until DECISION is received; /* remain blocked */
        write DECISION to local log;
    }
    if DECISION == GLOBAL_COMMIT
        write GLOBAL_COMMIT to local log;
    else if DECISION == GLOBAL_ABORT
        write GLOBAL_ABORT to local log;
} else {
    write VOTE_ABORT to local log;
    send VOTE_ABORT to coordinator;
}
```

# 2PC: Participant recovery (cont'd)

**Actions for handling decision requests:** /\* executed by separate thread \*/

```
while true {  
    wait until any incoming DECISION_REQUEST is received; /* remain blocked */  
    read most recently recorded STATE from the local log;  
    if STATE == GLOBAL_COMMIT  
        send GLOBAL_COMMIT to requesting participant;  
    else if STATE == INIT or STATE == GLOBAL_ABORT  
        send GLOBAL_ABORT to requesting participant;  
    else  
        skip; /* participant remains blocked */  
}
```

(b)

- Used to help other participants

# Next Time

- Byzantine Agreement and Recovery
- Read Chapter 8 TVS and FT\* paper